The `scidb` Package

Bryan W. Lewis
blewis@paradigm4.com

April 1, 2015

Contents

1 Introduction

SciDB is an open-source database that organizes data in n-dimensional arrays. SciDB features include ACID transactions, parallel processing, distributed storage, efficient sparse array storage, and native linear algebra operations. The `scidb` package for R provides two ways to interact with SciDB from R:

1. Through SciDB array classes for R. The arrays mimic standard R arrays and data frames, but operations on them are performed by the SciDB engine. Data are materialized to R only when requested.

2. By running SciDB queries from R, optionally returning results as data frames.

The SciDB array classes facilitate programming large-scale SciDB computation in R using natural R syntax. This vignette illustrates using SciDB from R by example. For more detailed information on the functions described in this vignette, see the manual pages in the package.

2 Connecting to SciDB

The `scidbconnect` function establishes a connection to a simple HTTP network service called shim running on a SciDB coordinator instance (see Section 12.3). The function may be safely called

multiple times. Once a connection is established, connection information is maintained until a different connection is established or the R session ends.

The network interface optionally supports SSL encryption and PAM or HTTP digest user authentication. The shim service can be configured to support either open/unencrypted or encrypted/authenticated ports, or both.

Connect to localhost by default on unencrypted port 8080:

```
> library("scidb")
> scidbconnect()
```

Connect to SciDB on an encrypted port 8083 with example PAM authentication:

```
> scidbconnect(host="localhost", port=8083,
+              username="scidbuser", password="test")
```

We recommend using only encrypted/authenticated sessions when communicating with SciDB over public networks.

2.1 Listing and removing SciDB arrays

The `scidblist` function lists SciDB objects (arrays, instances, etc.), optionally showing detailed schema information for arrays. Returned results may be filtered using regular expression-style syntax.

The `scidbremove` function removes a SciDB array, or optionally a set of SciDB arrays. The function accepts a vector of array names, resulting in the removal of all the specified arrays. Combine this feature with the regular expression filtering output of `scidblist` to remove sets of arrays matching the filter. Array names not associated with a specific R session are protected from easy removal to provide a modicum of protection for multiple SciDB users. Specify `force=TRUE` as indicated in the warning message to remove those arrays.

3   SciDB Arrays for R users

Data are organized by SciDB in n-dimensional sparse arrays. "Sparse" in SciDB arrays means that array elements may be left undefined, and such array elements are omitted from computations.

Note that this interpretation of sparse differs in a subtle way from that used by sparse matrices defined by R's Matrix package (whose sparse matrix elements are implicitly zero).

The elements of a SciDB array, called *cells*, contain one or more *attributes* (similar to R variables). The number and data types of attributes are uniform across all cells in an array, and SciDB stores data for each attribute separately. A one-dimensional SciDB array is conceptually similar to a data frame in R: the SciDB dimension index corresponds to data frame row index, and SciDB attributes to data frame columns. Higher-dimensional arrays in SciDB are generalizations of data frames and they don't correspond directly to objects in R; the `scidb` n-dimensional array class described below is sometimes limited to working with one attribute at a time to more closely emulate standard R arrays.

The `scidb` package defines two array classes for R with data backed by SciDB arrays: the `scidbdf` class for data frame-like objects, and the `scidb` class for matrix and general array objects.

3.1   The `scidbdf` data frame-like class

The `scidbdf` class defines a data frame-like class backed by one-dimensional SciDB arrays. Like data frames, the columns represent variables of distinct types and the rows represent observations. Each attribute in the backing SciDB array represents a column in the `scidbdf` object. The `scidbdf` object elements are read-only (the backing SciDB array may be manually updated, for example using the iquery function).

Use the `as.scidb` function to create new SciDB arrays and corresponding `scidbdf` R objects by copying R data frames into SciDB. The `types` and `nullable` options may be used to explicitly specify the SciDB type and nullability values of each data frame column. See the R help page for `as.scidb` and related `df2scidb` functions for more information.

The `scidb` function returns an R `scidbdf` or `scidb` object representation of an existing SciDB array.

Objects of class `scidbdf` obey a subset of R indexing operations. Columns may be selected by numeric positional index, variable name, or with the shorthand dollar sign notation similarly to standard data frames. Rows may be selected by numeric ranges–see the following section on array indexing for details.

Subsets of `scidbdf` objects are returned as new `scidbdf` objects of the appropriate size (dimension, number of attributes/columns). The package uses the special empty-bracket notation, `[]`, to indicate that data should be materialized to R as an R data frame. Illustrations are provided in the examples.

```
> # Upload the Michelson-Morley experiment data to SciDB,
> # returning a scidbdf data frame-like object:
> X <- as.scidb(morley, name="morely")
> str(X)


SciDB expression:  morely
SciDB schema:  <Expt:int32 NULL DEFAULT null,Run:int32 NULL DEFAULT null,
               Speed:int32 NULL DEFAULT null> [row=1:100,100,0]

Attributes:
  attribute  type nullable
1      Expt int32     TRUE
2       Run int32     TRUE
3     Speed int32     TRUE
Dimension:
  dimension start end chunk
1       row     1 100   100


> # Materialize the first four rows of X and the "Run" and "Speed"
> # columns to R (using [] to return results to R):
> X[1:4, c("Run","Speed")][]


  row Run Speed
1   0   1   850
2   1   2   740
3   2   3   900
4   3   4  1070
```

3.2 The scidb array class

Similarly to the data frame-like class, the package defines the scidb array class for R that represents vectors, matrices and general n-dimensional arrays. Array objects defined by the scidb class behave in some ways like standard R arrays. But their data reside in SciDB and most operations on them are computed by SciDB.

The scidb array class typically support working with a single array attribute at a time to conform to R arrays (which generally support a single value per cell). Note that, unlike R, SciDB array origin

indices are arbitrary and may be zero or negative (matrices and vectors in SciDB are generally zero-indexed). Data from `scidb` array objects are not materialized to R until extracted with the empty indexing function, `[]`. Additional notes and examples about indexing appear in the next section.

```
> # Upload the iris data to SciDB storing it in an array named `iris`
> # and create a data frame-like object `df` that refers to it:
> df <- as.scidb(iris)
> dim(df)


[1] 150    5


> # Alternatively, make a SciDB 1-d array object.
> x <- scidb(df, data.frame=FALSE)
> dim(x)


[1] 150


> # List all available attributes in the SciDB array:
> x@attributes


[1] "Sepal_Length" "Sepal_Width"  "Petal_Length" "Petal_Width"  "Species"
```

3.3 Arithmetic operations

The `scidb` array class supports a few standard linear algebra operations for dense and sparse matrices and vectors.

The example below shows a dense matrix example that compares matrix arithmetic in R and SciDB. Like previous data frame examples, use `as.scidb` to export R matrices and vectors to SciDB arrays.

We generally use `as.scidb` for convenience–it's far from the most efficient way to import data into SciDB. For very large data, use the SciDB bulk data load utility as outlined in the SciDB documentation instead.

```
> v <- as.scidb(matrix(rnorm(25),5))
> str(v)
```

```
SciDB expression:  R_array366d7949bc241102999947615
SciDB schema:  <val:double NULL DEFAULT null> [i=0:4,5,0,j=0:4,5,0]

Attributes:
  attribute   type nullable
1       val double     TRUE
Dimensions:
  dimension start end chunk
1         i     0   4     5
2         j     0   4     5


> crosprod( v[] )                    # Compute t(v) %*% v using R


             0             1             2             3             4
0  2.80783239 -0.05946079   0.6109525   1.53758978 -1.1351279
1 -0.05946079  0.37437931   0.4675420 -0.08090686   0.8351372
2  0.61095249  0.46754196   2.9929278 -2.32618189   0.7360878
3  1.53758978 -0.08090686  -2.3261819   5.48612039 -0.5832051
4 -1.13512788  0.83513722   0.7360878 -0.58320510   2.5291922


> # Now compute using SciDB, and materialize the result to R:
> crossprod(v)[]


             0             1             2             3             4
0  2.80783239 -0.05946079   0.6109525   1.53758978 -1.1351279
1 -0.05946079  0.37437931   0.4675420 -0.08090686   0.8351372
2  0.61095249  0.46754196   2.9929278 -2.32618189   0.7360878
3  1.53758978 -0.08090686  -2.3261819   5.48612039 -0.5832051
4 -1.13512788  0.83513722   0.7360878 -0.58320510   2.5291922
```

Basic matrix/vector arithmetic operations on SciDB arrays (addition, subtraction, matrix and matrix vector products, scalar products, `crossprod` and `tcrossprod`) use standard R syntax. You can mix R and SciDB matrices and vectors and the `scidb` package will try to do the right thing by assigning R data to temporary SciDB arrays conforming to required database schema. The next example shows an example of computations that mix `scidb` array objects with R vectors.

Expression	Operation	Operands	Output
A %*% B	Matrix multiplication	A, B Conformable SciDB arrays or R matrices/vectors	SciDB array
$A \pm B$	Matrix summation/difference	A, B SciDB arrays or R matrices/vectors	SciDB array
crossprod(A,B)	Cross product `t(A) %*% B`	A, B SciDB arrays or R matrices/vectors	SciDB array
tcrossprod(A,B)	Cross product `A %*% t(B)`	A, B SciDB arrays or R matrices/vectors	SciDB array
A */ B	Elementwise product/quotient	A, B Conformable SciDB arrays or R matrices/vectors	SciDB array
α */ A	Scalar multiplication/division	SciDB array A, scalar α	SciDB array
t(A)	Transpose	SciDB array A	SciDB array
sin(A)	Sine function, also other trig functions	SciDB array A	SciDB array
log($A, base$)	Logarithm function	SciDB array A, numeric base	SciDB array
diff($A, lag = 1$)	Finite differences	SciDB array A, integer lag	SciDB array
A[range, range, ...]	Subarray	SciDB array A	SciDB Array
A[ ]	Materialize	SciDB array A	R array
diag(A)	Matrix diagonal	SciDB array or vector A	SciDB arrays
svd(A)	Singular value decomposition	Dense SciDB array A	SciDB arrays
svd(A, nu)	Truncated SVD	Sparse or dense SciDB array A	SciDB arrays

Some scidb math operations

4   Subsetting and indexing SciDB objects

SciDB arrays are indexed by signed 62-bit integers. The origin of an array depends on its schema and can be any valid integer coordinate. However, zero is a typical origin, and some SciDB operations require zero-based indexes.

By contrast, R indices use positive 52-bit integer-valued double precision floating point numbers. That means that SciDB arrays can't always be indexed by R. In practice, R users should avoid creating SciDB arrays with coordinate systems that extend past R's limits. Moreover, all R array coordinate systems use 1-based indexing.

Subarray indexing operations use the SciDB convention. Thus, zero and negative indices are literally interpreted and passed to SciDB. In particular, negative indices do not indicate index omission, unlike standard R arrays.

Be aware that SciDB array coordinate systems may not begin at 1. Array coordinates are displayed

when data are downloaded, and the underlying coordinate system can be inspected using the `str` or `schema` functions.

Important indexing notes include:

- Use empty brackets, `[]`, to materialize data back to R. Otherwise, indexing operations produce new SciDB array objects.

- Use numeric indices in any dimension in the units of the underlying SciDB array coordinate system. Note that SciDB arrays generally are zero-indexed and may even have negative indices. SciDB data frame-like objects generally use 1-based indexing. When in doubt about the base index, use the `str`, `schema`, or `dimensions` functions to interrogate the SciDB object for details.

- Numeric indexing may include contiguous ranges or vectors of distinct coordinate values, but repeated coordinate values in a single dimension are not allowed. Examples of valid index ranges include `[1:4, c(3,1,5), -10:15]`, but not `[c(1,3,1)]`.

- The bracket subsetting function `[` takes an optional logical `redim` argument. When `FALSE`, the returned value has the same shape as the input value, with regions outside the indicated range marked sparse (empty). See the examples below.

- Subsetting operations using the bracket `[` function return zero-indexed results unless the `redim=FALSE` argument is specified.

- The `scidbdf` class represents 1-d SciDB arrays as data frame objects with SciDB array attribute as columns. Use either positional numeric or name-based indexing along columns, either with the dollar-sign notation or string indexing.

- All array types support labeled dimension indexing using the R `rownames`, `colnames`, or `dimnames` and related functions. Labels assigned in this way must be provided by 1-d SciDB arrays that map the integer coordinates to character label values. See the examples below.

- Arrays may be indexed by other SciDB arrays to achieve the effect of filtering by boolean expressions and similar operations, also illustrated below in the examples.

- Use the utility `between` function to avoid forming large sequences representing huge indexing ranges. For example, use `[between(1,1e9)]` instead of `[1:1e9]`. Use `between(0,Inf)`, for example, to indicate every index greater than or equal to zero.

- The `diag` function is supported to extract a vector of diagonal elements from a matrix, or to create a sparse diagonal matrix from a vector.

4.1   Indexing examples

Let's illustrate some indexing ideas by example. Consider the 5x4 matrix with unusual origin of $(-2, 0)$:

```
> (x <- as.scidb(matrix(as.double(1:20),nrow=5), start=c(-2,0)) )

A reference to a  5x4 SciDB array
```

The underlying coordinates are displayed when data are returned to R using the special empty index function:

```
> x[]

    0  1  2  3
-2 1  6 11 16
-1 2  7 12 17
0  3  8 13 18
1  4  9 14 19
2  5 10 15 20
```

The package adopts the convention that explicit subsets return zero-indexed results, noting also that zero and negative integers are valid SciDB indexes (deviating from R's usual interpretation):

```
> x[0:2, 2:3] []

   0  1
0 13 18
1 14 19
2 15 20
```

Use the optional `redim=FALSE` option to return a masked sparse array of the same shape as the input. This is often useful in conjunction with database join operations.

```
> x[0:2, 2:3, redim=FALSE] []
```

```
5 x 4 sparse Matrix of class "dgCMatrix"
    0 1  2  3
-2 . .  .  .
-1 . .  .  .
0  . . 13 18
1  . . 14 19
2  . . 15 20
```

The bracket function also supports the standard R `drop` argument:

```
> x[, 1] []

 0  1  2  3  4
 6  7  8  9 10

> x[, 1, drop=FALSE] []

   0
0  6
1  7
2  8
3  9
4 10
```

Non-contiguous selections are supported, noting again that returned subset results by convention here are zero-indexed:

```
> x[c(0,2,-1), ][]

   0  1  2  3
0 2  7 12 17
1 3  8 13 18
2 5 10 15 20
```

(However, unlike standard R array subsetting repeated indices are not supported.)

Coordinate axes may be labeled:

```
> rownames(x) <- letters[1:5]
> colnames(x) <- LETTERS[1:4]
> x[]
```

```
  A  B  C  D
a 1  6 11 16
b 2  7 12 17
c 3  8 13 18
d 4  9 14 19
e 5 10 15 20
```

The labels are maintained as ancillary SciDB arrays under the hood. You can view them:

```
> dimnames(x)
```

```
[[1]]
Reference to a SciDB vector of length 5
```

```
[[2]]
Reference to a SciDB vector of length Inf
```

```
> dimnames(x)[[1]] []
```

```
 -2  -1   0   1   2
"a" "b" "c" "d" "e"
```

If a sparse array has coordinates too large for R to manage, results are returned in an "unpacked" data frame form.

Note that SciDB objects with coordinates beyond the range of R display as inf. Coordinate labeling arrays are sparse with a very large upper coordinate bound by convention–allowing the number of labels to grow easily.

Labeled arrays can be subset, and the result retains labels:

```
> x[c('c','a','e'), c('B','A')] []
```

```
  A  B
a 1  6
c 3  8
e 5 10
```

Data-frame like objects are supported. Under the hood they are 1-d SciDB arrays.

```
> d <- as.scidb(iris)
> head(d)

  Sepal_Width Petal_Length Petal_Width Species
1         3.5          1.4         0.2  setosa
2         3.0          1.4         0.2  setosa
3         3.2          1.3         0.2  setosa
4         3.1          1.5         0.2  setosa
5         3.6          1.4         0.2  setosa
6         3.9          1.7         0.4  setosa
```

Note that the package will attempt to rename variables to conform to SciDB naming conventions when it has to.

The "columns" of a data frame like object are defined by the SciDB array attributes. You can use numeric or named indices, or the R dollar sign notation. The following subsets are all the same:

```
> d$Petal_Width
```

Reference to a SciDB vector of length 150

```
> d[, "Petal_Width"]
```

Reference to a SciDB vector of length 150

```
> d[, 4]
```

Reference to a SciDB vector of length 150

4.2   Comparison operators and filtering

The package follows R convention and returns a logical-valued object after comparison. Additionally, the package defines special comparison operators enclosed in the % symbol that return a sparse SciDB array object whose values not meeting the condition are masked as empty (preserving the values that do meet the condition).

The difference is illustrated in the following example.

```
> x <- as.scidb(head(iris))
> # Standard R-style logical comparison:
> (x$Petal_Length > 1.4 )[]

    1     2     3     4     5     6
FALSE FALSE FALSE  TRUE FALSE  TRUE

> # Special SciDB-style masked comparison:
> (x$Petal_Length %>% 1.4)[]

sparse vector (nnz/length = 2/6) of class "dsparseVector"
[1]    .    .    . 1.5    . 1.7
```

Comparison results can be used as filters in scidb and scidbdf objects. This works mostly like R, except that such filters may only be applied along coordinate axes. (That is, R's special single-index indexing mode is not supported.)

Additional examples showing dimension labels and filtering follow.

```
> # Example of labeling indices with a label array.
> set.seed(1)
> X <- as.scidb( matrix(rnorm(20),nrow=5) )
> # SciDB matrix objects like X default to zero-based indexing.
> # It's important that the label array have the same starting index:
> rownames(X) <- as.scidb( data.frame(letters[1:5]), start=0)
> X[c("b","a","d"), ]

A reference to a  3x4 SciDB array

> # Example of indexing by a condition on an auxiliary array.
> # Determine a subset of the rownames of X (another SciDB array):
> idx <- rownames(X) > "b"
> # Subset X by the condition on its rows:
> X[idx, ]

A reference to a  3x4 SciDB array
```

5 Data manipulation functions

The package defines a number of common SciDB data manipulation functions for SciDB array objects.

5.1 Apply, sweep and other array functions

R's `apply` function applies a function along margins of an array or matrix. A version of `apply` limited to SciDB aggregation functions is available for `scidb` and `scidbdf` objects. The SciDB aggregation function is supplied as a string and references the array attributes. The next example uses `apply` together with `sweep` to center the columns of a matrix.

```
> # Create an example matrix and upload to SciDB. Its single numeric array
> # attribute will be called "val" be default.
> A <- as.scidb( matrix(rnorm(100), nrow=10) )
> # Create a centered version of the matrix by computing its column means
> # with apply and subtracting them with sweep. Note that apply uses a
> # SciDB expression. We specify `eval=TRUE` below to compute and cache
> # the result of  sweep, improving the efficiency of the subsequent
> # crossprod operation.
>
> A0 <- sweep(A, 2, apply(A, 2, "avg(val)"), eval=TRUE)
> # The covariance matrix of A:
> crossprod(A0)/(nrow(A) - 1)

A reference to a  10x10 SciDB array
```

More general aggregation and data manipulation functions are outlined in the next section.

The `diff` function works similarly to R's usual one for computing finite differences along vectors and matrices.

Use the `cumulate` function to compute running operations along data, for example cumulative sums. The operation to be performed must be a valid SciDB aggregation function expressed as a character string. Here is a simple example:

```
> x <- as.scidb(iris)
> y <- cumulate(x, "sum(Petal_Width)")
> print(head(y, n=4))
```

```
   row Petal_Width_sum
1   1              0.2
2   2              0.4
3   3              0.6
4   4              0.8
```

The `count` function applied to a `scidb` array object returns the count of non-empty cells in in the backing SciDB array.

The `image` function displays a heatmap of a regrid of a 2-d `scidb` array object, and returns the regridded array to R. The `grid=c(m,n)` function parameter specifies the regrid window sizes in each array dimension, and defaults to the array chunk sizes. The regrid aggregation function may be specified using the `op` function argument, and by default averages the array values over the regrid windows.

5.2   Conditional subset operations

Use the `subset` function to filter array contents by a boolean expression somewhat similarly to the standard R `subset` function. Under the hood, this function uses the SciDB `filter` operator–the function name `subset` more closely matches standard R syntax.

The `subset` function requires two arguments, a SciDB array reference and a valid SciDB logical expression represented as a string. Here is a simple example:

```
> df <- as.scidb(iris)
> subset(df,"Petal_Width > 2.4")[]

    Sepal_Length Sepal_Width Petal_Length Petal_Width   Species
101          6.3         3.3          6.0         2.5 virginica
110          7.2         3.6          6.1         2.5 virginica
145          6.7         3.3          5.7         2.5 virginica
```

The `subset` function applied to arrays returns a sparse SciDB array with entries not meeting the filter condition omitted. This behavior differs from the standard R `subset` function, which can only operate on rows and columns. The `scidb` class also supports a short-hand notation for `subset` for infix comparison operations.

```
> # Upload a 5x5 matrix to SciDB. Its single numeric attribute name
> # defaults to 'val' (see help for as.scidb):
```

```
> set.seed(1)
> A <- as.scidb( matrix(rnorm(9),nrow=3) )
> # Apply a filter to the SciDB array and return the results to R:
> subset(A, "val>0")[]

3 x 3 sparse Matrix of class "dgCMatrix"
          0         1         2
0 .         1.5952808 0.4874291
1 0.1836433 0.3295078 0.7383247
2 .         .         0.5757814

> # Alternatively, achieve the same filter with an infix comparison:
> (A > 0)[]


      0     1    2
0 FALSE  TRUE TRUE
1  TRUE  TRUE TRUE
2 FALSE FALSE TRUE
```

5.3 Database joins

The package supports a limited form of the R merge function, enabling a number of database join-like operations on SciDB array objects. See help("merge",package="scidb") function for detailed help. Here is an example that performs an inner join on array attributes:

```
> authors <-
+   data.frame(
+     surname = c("Tukey", "Venables", "Tierney", "Ripley", "McNeil"),
+     nationality = c("US", "Australia", "US", "UK", "Australia"),
+     deceased = c("yes", rep("no", 4)),
+     stringsAsFactors=FALSE)
> books <-
+   data.frame(
+     name = c("Tukey", "Venables", "Tierney",
+              "Ripley", "Ripley", "McNeil", "R Core"),
+     title = c("Exploratory Data Analysis",
+               "Modern Applied Statistics ...",
+               "LISP-STAT", "Spatial Statistics", "Stochastic Simulation",
```

```
+                 "Interactive Data Analysis", "An Introduction to R"),
+      other.author = c(NA, "Ripley", NA, NA, NA, NA, "Venables & Smith"),
+      stringsAsFactors=FALSE)
> a <- as.scidb(authors)
> b <- as.scidb(books)
> merge(a,b,by.x="surname",by.y="name")[,4:8][]

  row_2 nationality deceased    name                          title
1     0   Australia       no  McNeil    Interactive Data Analysis
2     1          UK       no  Ripley            Spatial Statistics
3     2          UK       no  Ripley          Stochastic Simulation
4     3          US       no Tierney                     LISP-STAT
5     4          US      yes   Tukey    Exploratory Data Analysis
6     5   Australia       no Venables Modern Applied Statistics ...
  other_author
1         <NA>
2         <NA>
3         <NA>
4         <NA>
5         <NA>
6       Ripley

> # cf. The standard R data.frame merge function:
> # merge(authors,books,by.x="surname",by.y="name")
```

The `merge` implementation has some limitations outlined in its man page. Joins on attributes (instead of along coordinate axes) are presently the most limited kinds and outer joins are only supported in certain cases.

5.4   Aggregation

The `aggregate` function performs various aggregation operations on SciDB array objects. Aggregation may be defined over array dimensions and/or array attributes very similarly to standard R aggregation syntax, except that aggregation functions must generally be valid SciDB aggregate expressions, represented as strings.

```
> # Aggregation example
> df <- as.scidb(iris)
> aggregate(df, by="Species", FUN="avg(Petal_Length), stdev(Petal_Width)")[]
```

```
   Petal_Length_avg Petal_Width_stdev    Species
0             1.462         0.1053856     setosa
1             4.260         0.1977527 versicolor
2             5.552         0.2746501  virginica
```

Although limited support for aggregation by data frame column is available (see the example in this section), the `aggregate` function performs best along SciDB array dimensions. The `aggregate` function also supports single- and multi-dimensional moving window aggregation along array coordinate systems, and 1-d aggregation along consecutive data values in sparse arrays. See the on-line man page for additional details and examples.

Alternatively, use the SciDB-specific `redimension` function to perform aggregations. This form of aggregation is sometimes more efficient in the database. See the help page for `redimension` and the examples in the following section on `redimension`.

5.5 Binding new variables to an array

The SciDB package defines the `bind` function to add variables to arrays similar to the R `cbind` function for data frames. However, `bind` can also operate on higher-dimensional arrays. The example below adds a variable named 'prod' to the SciDB array `df` defined in the last example.

```
> y  <- bind(df, "prod", "Petal_Length * Petal_Width")
> head(y, n=3)

  Sepal_Width Petal_Length Petal_Width Species prod
1         3.5          1.4         0.2  setosa 0.28
2         3.0          1.4         0.2  setosa 0.28
3         3.2          1.3         0.2  setosa 0.26
```

5.6 SciDB redimension

The `redimension` function is a wrapper for the SciDB `redimension` operator–the Swiss army knife of SciDB. The `redimension` function can be used to change the dimensionality of an array and even to apply reduction functions while doing so to achieve aggregations.

The following example illustrates using `redimension` to compute a congingency table along levels of two variables, "class" and "Species."

```
> x <- as.scidb(iris)
> # bind an example new 'class' column:
> y <- bind(x, "class", "iif(Petal_Width>2, int64(1), 0)")
> # Contingency table along two dimensions:
> redimension(y, dim=c("class", "Species"), FUN=count) []

2 x 3 sparse Matrix of class "dgCMatrix"
  setosa versicolor virginica
0    50         50        27
1     .          .        23

> # Aggregation example (cf. aggregation example above)
> redimension(x,
+             dim="Species",
+             FUN="avg(Petal_Length) as Petal_Length_avg,
+                 stdev(Petal_Width) as Petal_Width_stdev")[]

           Petal_Length_avg Petal_Width_stdev
setosa                1.462         0.1053856
versicolor            4.260         0.1977527
virginica             5.552         0.2746501
```

The function creates and applies row and dimension names as required.

5.7   Sorting and enumerating factors

Use the `sort` function to sort on a subset of dimensions and/or attribute of a SciDB array object, creating a new sorted array.

Use the `unique` function to return a SciDB array that removes duplicate elements of a single-attribute SciDB input array.

Use the `index_lookup` function along with `unique` to bind a new variable that enumerates unique values of a variable similarly to the R `factor` function.

Use the new `factor_scidb` function to create R factor variables whose levels are determined by a SciDB dimension array.

Examples follow. Note that we use the SciDB `project` function in one example. `project` presents an alternative syntax to the functionally equivalent column subset selection of variables using brackets.

```
> x <- as.scidb(iris)
> # Sort x by Petal_Width and Species
> a <- sort(x, attributes=c("Petal_Width","Species"))
> head(a, n=3)

  Sepal_Width Petal_Length Petal_Width Species
1         3.1          1.5         0.1  setosa
2         3.0          1.4         0.1  setosa
3         3.0          1.1         0.1  setosa

> # Find unique values of Species:
> unique(x$Species)[]

      Species
0      setosa
1 versicolor
2  virginica

> # Add a new variable that enumerates factor levels of Species:
> head(index_lookup(x, unique(x$Species), "Species"))

  Sepal_Width Petal_Length Petal_Width Species Species_index
1         3.5          1.4         0.2  setosa             0
2         3.0          1.4         0.2  setosa             0
3         3.2          1.3         0.2  setosa             0
4         3.1          1.5         0.2  setosa             0
5         3.6          1.4         0.2  setosa             0
6         3.9          1.7         0.4  setosa             0
```

See the on-line documentation for `factor_scid` for additional examples.

5.8 Missing values

SciDB attribute values within a cell may be explicitly marked missing, indicated by a special SciDB missing code also referred to as a `NULL` code. SciDB internally supports a large number of possible missing codes. All SciDB missing code values are mapped to `NA` values in R.

6 Lazy evaluation, array promises, and garbage collection

Most operations on SciDB array objects return array promises–new SciDB array objects that have not been fully evaluated by SciDB yet, but that promise to return values when asked. SciDB array promises are essentially just SciDB query expressions together with a schema that the resulting output array will have once it's been evaluated and an environment providing context for the evaluation. In such cases, the `name` slot of a `scidb` or `scidbdf` object shows the SciDB query expression.

Sometimes it can be more efficient to explicitly evaluate and cache a SciDB array then to use it as an array promise. This can always be achieved by using the `scidbeval` function, or by using the `` `eval` ``=TRUE argument in the functions that support that. Dynamically allocated arrays use a naming convention that begins with "R_array" and end with a unique numeric identifier determined by the current SciDB session.

Ephemeral intermediate arrays are by default connected to R's garbage collector and automatically deleted from the SciDB catalog when they are no longer needed by R. Users can disconnect SciDB array objects from R's garbage collector (making stored arrays persistent in SciDB), by using the `gc=FALSE` flag in the `as.scidb` and `scidbeval` functions.

SciDB arrays keep references to other SciB array dependencies in a list within an environment, preventing them from automatic garbage collection so that the promise can be evaluated. The environment is slot is named `@gc` in the array objects and the dependency list is `@gc$depend`.

Here is an example:

```
> # Create a 5x4 SciDB matrix named 'test' in the database. We set `gc=TRUE`
> # so that the SciDB array `test` will be deleted by R's garbage collector
> # for us automatically:
> A <- as.scidb(matrix(rnorm(20),nrow=5), name="test", gc=TRUE)
> A@name

[1] "test"

> # Assign the third row of the transpose to a scidb object B:
> B <- t(A)[3,]
> # B is a promise object--it's name is an (unevaluated) SciDB query expression
> # that depends on A. But it has a valid output schema, and we also see that
> # it depends ultimately on the SciDB array `test`:
> B@name
```

```
[1] "redimension(
      subarray(
        redimension(
          between(
            transpose(test),3,null,3,null),
          <val:double NULL>[j=3:3,4,0,i=0:4,5,0]),
        null,null,null,null),
      <val:double NULL>[i=0:4,5,0])"
```

```
> B@schema
```

```
[1] "<val:double NULL DEFAULT null> [i=0:4,5,0]"
```

```
> B@gc$depend[[1]]@gc$depend[[1]]@name
```

```
[1] "transpose(test)"
```

```
> # We can use `scidbeval` to force B's evaluation and storage inside SciDB:
> C <- scidbeval(B)
> # C has the same schema as B, but is now evaluated and stored inside SciDB
> # using an automatically-generated name. Unlike B, C does not depend on
> # the `test` array any more.
> C@name
```

```
[1] "R_array366d7fc6adf21102999947615"
```

```
> C@schema
```

```
[1] "<val:double NULL DEFAULT null> [i=0:4,5,0]"
```

7   SciDB-specific functions

This section briefly summarizes a number of utility functions that are simple wrappers of miscellaneous SciDB operators. Except for count, which returns a number, the listed functions return references to new SciDB array objects.

– attribute_rename Rename one or more attributes of a SciDB array.

- – build Creates new dense arrays from SciDB expressions and schema.
- – cast Create a new array by altering the schema of an array.
- – count Return the number of nonempty cells of an array.
- – cumulate Compute a cumulative aggregate function, for example a cumulative sum or product.
- – dimension_rename Rename one or more dimensions of a SciDB array.
- – index_lookup Look up attribute values in a mapping array, adding a new attribute with the index.
- – project Pick out one or more attributes from an array.
- – redimension Redimension an array into the provided new schema.
- – regrid Decimate an array by binning it along one or more dimensions and applying an aggregate function to the bins.
- – rename Rename a SciDB array.
- – repart Repartition an array.
- – reshape Reshape an array.
- – schema Show the schema of a SciDB array object.
- – scidblist List SciDB arrays (also scidbls).
- – scidbremove Delete SciDB arrays (also scidbrm).
- – slice Slice a SciDB array along one of its axes, returning a lower-dimensional array.
- – sort Sort the elements of an array.
- – subarray Perform a SciDB subarray or between operation on an array.
- – unique Return a set of unique elements of a sorted array.
- – unpack Unpack a SciDB array into a 1-d table.
- – xgrid Prolong one or more coordinate axes by replicating cell values.

See the online manual pages for more details and examples.

8   Miscellaneous functions

This catch-all section summarizes miscellaneous functions.

8.1 Histograms

The generic function `hist` computes a histogram of the given data values in a SciDb array object. If `plot=TRUE`, the resulting object of class histogram is plotted by `plot.histogram`, before it is returned.

This histogram function only supports equidistant breaks. If `right=TRUE` (default), the histogram cells are intervals of the form '(a, b]', i.e., they include their right-hand endpoint, but not their left one.

8.2 Last-value imputation

The `na.locf` function is a generic function for replacing each missing or empty value with the most recent non-missing value prior to it.

Unlike the usual `na.locf` function from the `zoo` package, the SciDB `na.locf` function fills in both missing (SciDB `null` values) and empty (SciDB sparse) values with the last non-missing and non-sparse value along the indicated dimension.

Time series represented in SciDB are often sparse arrays. The `na.locf` function defines a conveniet way to fill in all missing values along the time coordinate axis down to the time resolution.

Caution! The output array is a mostly dense, filled-in version of the input array. If the time resolution is very fine and the input array very sparse, then the output array can be huge. Consider using `regrid` first on very fine time scales to reduce their resolution.

The default `na.locf` method in the `zoo` package unforunately overrides this function (it uses `ANY` in its method signature). If you need to use SciDB arrays and the `zoo` package, prefix SciDB's version with `scidb::na.locf(...)`.

9 Running SciDB queries

The `iquery` function executes SciDB queries using either the SciDB array functional language (AFL) or declarative array query language (AQL) syntax. When AFL is used, the `iquery` function optionally returns query results in an R data frame if the argument `return=TRUE` is specified. Returned output is similar to output obtained by the SciDB `iquery` command-line program with the `-olcsv+` option. The `iquery` function does not return anything by default.

Query results returned by the `iquery` function are internally presented to R using a generic CSV format, providing very flexible support for many data types. (The *n*-dimensional array class described in the next section uses a binary data exchange method between R and SciDB.) Note that, although R and SciDB have a number of common data types, each system contains types not supported by the other. Thus, conversion errors may arise. The `iquery` function is designed to reasonably minimize such issues and simplify basic data transfer between the systems. Data types common to R and SciDB include double-precision numeric, character string, logical, and 32-bit integers. The `iquery` function supports standard R `read.table` parameter options to facilitate type conversion.

The following code example illustrates basic use of `iquery`.

```
> # Manually build a 1-d SciDB array named "P:"
> iquery("store(build(<x:double>[i=0:99,100,0],asin(1)*i/25),P)")
> # Return a query as a data.frame, manually specifying returned column classes:
> S <- iquery("apply(P,y,sin(x))",return=TRUE,
+             colClasses=c("integer","double","complex"))
> head(S, n=3)

  i          x             y
1 0 0.00000000 0.00000000+0i
2 1 0.06283185 0.06279052+0i
3 2 0.12566371 0.12533323+0i
```

9.1 Iterating over query results

The `iquery` function returns query results into a single R data frame by default. Large results expected to contain lots of rows may be iterated over by setting the `iterative=TRUE` argument. When `iquery` is used with the `iterative=TRUE` setting, it returns an iterator that iterates over chunks of rows of the result data frame. Iterators are defined by the `iterators` package. Their data may be directly accessed with the `nextElem` method, or indirectly with `foreach`. See the `iterators` and `foreach` packages for many examples and further documentation of their use.

```
> # Build a small 1-d SciDB test array:
> iquery("store(build(<x:double>[i=1:10,10,0],i/10.0),A)")
> # Return the result of a SciDB apply operator in an R iterator with a
> # block size of at most 7 rows at a time:
```

```
> library("iterators")
> it <- iquery("apply(A,y,sqrt(x))", return=TRUE, iterative=TRUE, n=7)
> nextElem(it)

  i   x          y
1 1 0.1 0.3162278
2 2 0.2 0.4472136
3 3 0.3 0.5477226
4 4 0.4 0.6324555
5 5 0.5 0.7071068
6 6 0.6 0.7745967

> nextElem(it)

   i   x          y
1  7 0.7 0.8366600
2  8 0.8 0.8944272
3  9 0.9 0.9486833
4 10 1.0 1.0000000
```

10   Error handling

SciDB errors are trapped and converted to R errors that can be handled by standard R mechanisms. Some operations might try to return too much data to R, exceeding R's indexing limitations, system memory, or both. The package tries to avoid this kind of error using package options that limit returned data size shown in the package options section below.

11   Function and method list

Here is a summary list of functions and methods defined by the package. See online help for examples and more information on each one.

```
<                        <=                       ==
>                        >=                       -
!=                       /                        [
```

[<-	$	*
%<=%	%<%	%==%
%>=%	%>%	%*%
+	abs	acos
aggregate	all.equal	antijoin
apply	asin	as.scidb
atan	attribute_rename	bernoulli
between	bind	bound
build	c	cast
cbind	chunk_map	colnames
colnames<-	cos	count
crossprod	cumulate	default
df2scidb	diag	diff
dim	dimension_rename	dimensions
dimnames	dimnames<-	dist
exp	factor_scidb	Filter
glm	glm.fit	glm_scidb
head	hist	image
index_lookup	iqdf	iquery
is.scidb	is.scidbdf	is.temp
kmeans	lag	length
levels_scidb	log	max
mean	median	merge
min	model_scidb	na.locf
names	names<-	ncol
nrow	Ops	order
peek	persist	phyper
predict	print	project
qhyper	quantile	rank
rbind	redimension	regrid
remove_old_versions	rename	repart
replaceNA	reshape	rownames
row.names	row.names<-	rownames<-
schema	scidb	scidb_attributes
scidbconnect	scidb_coordinate_bounds	scidb_coordinate_chunks
scidb_coordinate_end	scidb_coordinate_overla	scidb_coordinate_start
scidbdf	scidbdisconnect	scidbeval
scidb_fisher.test	scidblist	scidbls

scidb_nullable	scidbremove	scidbrm
scidb_types	sd	show
show_commit_log	sin	slice
solve	sort	sqrt
str	subarray	subset
sum	summary	svd
sweep	t	tail
tan	tcrossprod	tsvd
unbound	unique	unpack
var	xgrid	

12   Package installation

Installation proceeds in two steps: installing the R package on any computer that has a network connection to a SciDB database, and installing a simple network service on the SciDB database coordinator computer.

12.1   Installing the R package from sources on GitHub

The scidb package source is maintained in the SciDBR GitHub repository. That's where the most up-to-date version of the package is available. Released versions of the package posted to CRAN are updated much less frequently, approximately semiannually. A git tag indicates each CRAN release version of the package in the source code repository.

The wonderful devtools R package makes installation of source packages from GitHub nearly as simple as installation from CRAN.

```
library("devtools")
install_github("SciDBR", "Paradigm4", quick=TRUE)
```

12.2   Installing the R package from CRAN

The scidb package is available on CRAN. Start an R session and run:

```
install.packages("scidb")
```

12.3   Installing the simple HTTP service for SciDB

The SciDB R package requires installation of a simple open-source HTTP network service called `shim` on the computer that SciDB is installed on. The service needs to be installed only on the SciDB coordinator computer, not on client computers that connect to SciDB from R. It's available in packaged binary form for supported SciDB operating systems, and as source code which can be compiled and deployed on any SciDB installation.

Both installation approaches install the `shim` network service on the SciDB coordinator computer. Installing as a service requires root permission. The compiled source code version requires no special permissions to run.

Installation from binary software packages for SciDB-supported operating systems is easiest. Detailed up-to-date information can be found on Paradigm4's laboratory on Paradigm4's Github repository at https://github.com/Paradigm4/shim/wiki/Installing-shim.

The installed `shim` network service exposes SciDB as a very simple HTTP API. It includes a simple browser-based status and query tool. After installing shim, point your browser to the I.P. address of the SciDB coordinator machine and port 8080. Note that this API is not official and may change in the future. Help drive those changes by contributing ideas, code and bugfixes to the project on Github, or feel free to discuss the service on the SciDB.org/forum or via Github issues.

12.4   Package options, miscellaneous notes, and software license

The `scidb` package defines several global package options. Package options may be set and retrieved with the R `options` function, and are listed in Table 2.

Miscellaneous notes:

- R does not support 64-bit integer types. 64-bit signed and unsigned integers smaller than 2^{53} in magnitude will be represented as double-precision floating point numbers. 64-bit integers outside that range appear as $+/-$`Inf`. All other integers (int8, uint8, int16, uint16, etc.) are represented in R by 32-bit signed integers. The SciDB uint32 type is not supported.

- R doesn't support single-precision floating point numbers. `iquery` results convert single-precision numbers within SciDB to double-precision floating-point numbers in R. Single-precision SciDB numbers are not supported by the `scidb` array class.

- SciDB does not natively support complex numbers. Loading complex numbers directly into SciDB from R is not defined.

```
# Maximum allowed sequential index limit (for larger, use between)
  options(scidb.index.sequence.limit=1000000)
# Maximum allowed elements in an array return result
  options(scidb.max.array.elements=100000000)
# The scidb.version option is set during scidbconnect(). However, users
# may carefully override it to enable certain bug fixes specific to older
# versions of SciDB.
  options(scidb.version=13.9)
# Set this to 32 for SciDB version 13.6!
  options(scidb.gemm\_chunk_size=1000)
# Default shim port and host.
  options(scidb.default_shim_port=8080L)
  options(scidb.default_shim_host="localhost")
# Make it harder to remove arrays. When this option is TRUE, users
# have to specify scidbrm(array, force=TRUE) to remove arrays that do not
# begin with "Rarray".
  options(scidb.safe\_remove=TRUE)
# Disable SSL certificate host name checking by default. This is important mostly
# for Amazon EC2 where hostnames rarely match their DNS names. If you enable this
# then the shim SSL certificate CN entry *must* match the server host name for the
# encrypted session to work. Set this TRUE for stronger security (help avoid MTM)
# in SSL connections.
  options(scidb.verifyhost=FALSE)
# Set to FALSE to disable user-interruptable HTTP transactions.
  options(scidb.interrupt=TRUE)
# Set to TRUE to enable experimental shim stream protocol, avoids copying query
# output to data file on server # (see https://github.com/Paradigm4/shim).
  options(scidb.stream=FALSE)
```

Package options

- The `iquery` function provides the most flexible mechanism for type conversion between the systems, fully under user control using `read.table` options.

- Naming convention and associated restrictions vary between R and SciDB. For example, SciDB does not allow decimal points in attribute names. The package may alter names with character substitution to reconcile names when it is reasonable to do so. A warning is emitted whenever an object is automatically renamed in this way.